

COURTFLOW-GS: Online Multi-View Dynamic Gaussian Reconstruction with Fused MLS Tracking

Shangar Muhunthan

September 28, 2026

Abstract

COURTFLOW-GS turns synchronized videos from cameras around a basketball court into a 3D scene that can be viewed from a new camera angle. It first builds one clean, static “starting scene” from frame 0. It then follows each player through the video by moving small sets of control points, and occasionally improves visual detail. The result is a full 3D reconstruction, novel-view videos, and a saved Gaussian model that can be rendered later without retraining.

1 Problem and data

The data contains 36 synchronized 4K cameras placed in a ring around the court. Their positions, orientations, and lenses have already been calibrated, so the system knows exactly where every camera is. We train with every third camera (views $0, 3, \dots, 33$) and keep the remaining 24 cameras aside for an honest test from viewpoints the model did not see during training. The camera calibration is fixed throughout the method.

The input adapter reads the 36 videos and the COLMAP calibration (`cameras.txt` and `images.txt`) and creates an undistorted working copy for training. By default this copy is half resolution (1920×1080), which keeps the camera geometry while making the many repeated training renders cheap; the same pipeline also runs at the native 3840×2160 .

2 Existing methods and the 12-view problem

What existing dynamic methods assume. Recent dynamic Gaussian methods reconstruct multi-view video either one frame at a time (*frame-stream*) or in short clips. Dynamic 3D Gaussians [1] gives every Gaussian its own position and rotation at every timestep and regularizes them with local-rigidity losses. 3DGStream [2] learns a small neural transformation cache per frame; QUEEN [3] learns quantized per-Gaussian attribute residuals per frame. ClipGStream [4] optimizes short clips with spatio-temporal fields and inherits anchors from clip to clip. SC-GS [5] is the closest in its motion model: sparse control points drive dense Gaussians, but their motion is predicted by a deformation MLP trained offline over the whole sequence.

These methods are evaluated with many more cameras than this assignment allows (Table 1). The closest benchmark is ClipGStream’s Long 360: a 36-camera 4K ring like ours, but trained on 32 cameras and tested on 4. Here we train on 12 and test on 24. With cameras 30° apart, a small player seen by a few cameras is only weakly constrained. Per-Gaussian motion or per-Gaussian residuals then have enough freedom to fit the training images without being correct from new viewpoints.

Table 1: Published settings of related methods. Numbers are from the papers, on their own datasets and GPUs, so they show the regime each method was designed for, not a like-for-like comparison.

Method	Training cameras (bench-mark)	Motion model	Reported training time
Dynamic 3D Gaussians	27 of 31 (Panoptic dome)	per-Gaussian 6-DOF per frame	2 h for 150 frames, RTX 3090
3DGStream	20 of 21 (N3DV); 13 of 14 (Meet Room)	neural transform cache per frame	12 s/frame + 23 min first frame, RTX 3090
QUEEN	20 of 21 (N3DV)	quantized per-Gaussian residuals	6 s/frame, A100
ClipGStream	32 of 36 (Long 360, 4K ring)	per-clip spatio-temporal fields	≈ 0.5 h per sequence
SC-GS	dense / monocular	control points + deformation MLP	offline, whole sequence
COURTFLOW-GS (ours)	12 of 36 (4K ring)	per-player control-point MLS	Table 6

ClipGStream on this dataset. Before COURTFLOW-GS, we ran ClipGStream on exactly this capture with the same 12 training cameras and 24 held-out cameras (RTX 3080, 960×540 , first 10-frame clip). Table 2 shows the sparse-view failure directly: training-view PSNR is 26–28 dB, but held-out PSNR is only 17.2–17.9 dB. Training longer made the model *worse* on new views (17.88 dB after 6000 iterations, 17.53 dB after 40000) while training-view PSNR rose. Its reference clip took 7.2 min and each following 10-frame clip took 132–143 s (about 13.5 s per frame on the 3080).

Table 2: ClipGStream on this capture (12 training, 24 held-out cameras, RTX 3080). Mean over the 24 held-out cameras at 960×540 .

Run	Train PSNR	Held-out PSNR	Held-out SSIM	Held-out LPIPS
6000 iterations	25.96	17.88	0.543	0.308
40000 iterations	28.22	17.53	0.522	0.320
40000 iterations, court-focused	–	17.22	0.498	0.342
Streamed sequence, first clip	–	17.93	0.546	0.306

How COURTFLOW-GS differs. COURTFLOW-GS is designed around the sparse ring rather than adapted to it:

- **Few motion unknowns.** Players move through 1024 control points each, with one shared translation per player and an as-rigid-as-possible shape term, instead of free per-Gaussian motion. Twelve views can constrain a few thousand control offsets; they cannot constrain millions of independent Gaussian updates.
- **Geometry from masks and matching, not only photometry.** Sparse-view 3DGS overfits because colour alone admits many wrong geometries. COURTFLOW-GS seeds players from occlusion-aware hulls and masked MAST3R crop matches, and uses DropGaussian [6]-style regularization, so held-out views start from a plausible 3D shape.
- **Identity is solved in 3D.** Every player is its own motion group, so a teammate or the court cannot absorb another player’s motion, and SAM2 identity swaps are repaired in 3D.
- **Static work is done once.** About 95% of each image is static court and stands. Frame-stream methods re-fit that area every frame; COURTFLOW-GS freezes it after frame 0 and spends per-frame computation only on the moving players. This is the rare case where the same change improves both speed and quality: less compute, and no drift in the static scene.

The one-shot geometry work costs more at frame 0 (Stage A) than a plain 3DGS initialization, but it is paid once per scene; afterwards each frame costs a few seconds.

3 Methodology

Before doing anything else, every image is undistorted. This removes lens bending, so a straight camera ray in the mathematics corresponds to a straight ray in the image. In the notation used later, camera c is $P_c = K_c[R_c | t_c]$, with world-to-camera conversion $x_c = R_c x + t_c$. Readers do not need this equation to follow the method: it only says that the known calibration projects a 3D point into each image. Camera poses, lens parameters, and per-camera brightness corrections are kept fixed.

3.1 Frame-0 Reconstruction

Masks and 3D identities: Photometric reconstruction alone is too ambiguous for the moving players. Each player occupies a small portion of a frame, shares colour and texture with teammates, and is frequently overlapped by another player or various structures / elements in the scene. Foreground masks tell the system both “which pixels belong to this person for geometry matching” and “where this person is allowed to appear when rendered.” This prevents background splats from being absorbed into players and prevents players from growing fuzzy opaque halos onto the court. It also tells the optimizer which Gaussian/control-point group should explain each pixel. This prevents a global reconstruction from improving RGB error by assigning one player’s appearance or motion to another player or to the background.

Simply running a detector in each camera is not enough: “person 4” in one image has no reason to be the same physical person as “person 4” in another. We first tried matching 2D mask centres and colours between cameras. With cameras about 30° apart, that split 68 masks into 60 supposed people. Full-frame MAST3R matching was also too coarse for small players and made unreliable star-shaped rays. Instead, CourtFlow-GS establishes identity *in 3D first*.

Faster R-CNN proposes person boxes (score ≥ 0.5), and SAM2 turns each box into a precise foreground mask. We use agreement between those masks to make rough 3D blobs, where each blob is a likely physical player. We project each blob back into every camera and match it to the local SAM2 mask. When a mask is missing, the projected blob tells SAM2 where to look again. In this way, noisy per-camera masks become one consistent player ID across the ring.

Concretely, we cover the court with a 2.4 m-high box that extends 1 m beyond the court lines, and divide it into small cubes called voxels. A voxel is kept only if enough cameras see it inside a person mask: at least $\max(4, n_{\text{visible}} - 2)$ cameras. Put simply, several cameras must agree, but up to two are allowed to disagree because a player can be hidden. We discard blobs that are too small, too short, or floating above the floor. A blob wide enough to contain two people is split on the court floor.

For a missing camera mask, we accept the new SAM2 result only if its colour is similar to that same 3D player in the other cameras. If it still looks wrong, we mark that player as *unknown* in that camera, not as background. This is important: a hidden or missed player should not teach the model that the court is visible there. Finally, if two accepted masks overlap, the closer player owns the pixel. The result is one label image L_c per camera in which every pixel is either background or exactly one player. Later, this lets every player have separate Gaussian splats and separate motion controls.

Occlusion-aware visual hull. A visual hull is a rough 3D body made by intersecting the player silhouettes from several cameras. The usual rule says to remove a voxel if even one camera does not see the player there. That works for a clean object on a turntable, but fails in basketball: another player, the hoop support, or padding can hide a real arm, leg, or head. We saw those missing body parts with a strict hull.

The opposite rule (forgiving any missing view) was also bad. One allowed miss made player hulls about $2.5\times$ too large. Our compromise asks one simple question: *is there measured evidence that something closer to this camera was blocking the player?* We forgive a missing pixel only when the answer is yes. This is what “occlusion-aware” means here.

For each player, we test a fine 1 cm grid inside the rough 3D blob, with a small margin. Camera c votes against a candidate voxel x only when the voxel is visible in that camera, the player has a trusted label there, the projected pixel is not labelled as that player, and nothing nearer explains why. A nearer explanation can be another labelled player or a measured background surface (such as the hoop support) at least 30 cm in front. The background depth map D_c comes from reliable full-frame MAST3R background points that agree in at least two additional cameras. We retain the nearest depth at each image location and ignore floor points within 0.5 m of the court.

We keep x when at least two cameras see it as the player and no camera has an *unexplained* reason to reject it. Voxels below the floor are then removed. A camera where the player is unknown casts no vote. Thus the method does not keep geometry merely because occlusion might be possible; it keeps it only when a nearer player or surface actually supports that explanation.

Masked crop geometry and Gaussian seeds. Feature matching across a full image works for the court, but a distant player may occupy only about 6×25 pixels once MAST3R shrinks the whole frame to 512 pixels. Our full-frame player matches made a radial “star” of unreliable 3D rays. The masks solve this by giving the player the full matching resolution: for nearby camera pairs, we crop a padded square around the same 3D player and resize that crop to 512×512 . MAST3R searches only inside the source player mask.

After mapping matches back into the original camera images, we retain only confident matches that are in front of the cameras, reproject accurately, use cameras far enough apart for useful depth, lie inside a slightly expanded player hull, and agree with at least one more labelled camera. This removes the court and background rays that slipped through earlier mask-only tests.

The retained matches are accurate surface seeds. Hull voxels fill only textureless gaps more than two hull voxels away from a seed. We colour points only from cameras that agree with the label and can actually see the point, remove outliers, merge nearby player points at 5 mm and background points at 10 cm, and turn the result into oriented, elongated Gaussian splats. Each Gaussian stores position, shape, opacity, and simple view-dependent colour (degree-1 spherical harmonics). Players train for 7000 iterations (four cameras per step) and background for 10000 (two cameras per step), using

$$\mathcal{L}_{\text{image}} = \|I - \hat{I}\|_1 + 0.2\text{D-SSIM} + 0.1\mathcal{L}_{\text{alpha}} + 0.05\mathcal{L}_{\text{depth}}.$$

This loss asks the rendered image to match the real image, its structure, its silhouette, and its initial depth. The depth term fades out halfway through training, after it has helped establish the coarse shape. Starting at iteration 500, we add detail where needed every 100 iterations until halfway through training. We remove nearly transparent splats and splats that land outside an expanded player mask in two cameras.

3.2 Group-aware control-point MLS

Control Groups After frame 0, the task is to move the same Gaussian splats as players move. Rather than moving millions of splats independently, we place a much smaller set of control points on each player. Moving a control point moves nearby splats smoothly. This is called moving-least-squares (MLS) deformation.

A single global control graph is reasonable for one connected object, but unsafe in a crowded game. Controls on two nearby players can otherwise explain the same image error, causing one player to pull another player’s splats or the background. CourtFlow-GS treats every identity as a separate motion group. This keeps the useful smooth motion of MLS while forbidding cross-player influence.

Every dynamic Gaussian belongs to exactly one person; the ball, when enabled, is another group, and background is group 0. We choose controls spread evenly over each object: 1024 per person, 16 for the ball,

and 1000 for background. Each Gaussian is permanently connected to its eight nearest controls from the *same* group. Closer controls get more influence through the fixed radial weights

$$w_{ij} = \frac{\exp(-\|x_i - p_j\|^2 / (2\sigma_j^2))}{\sum_{\ell=1}^8 \exp(-\|x_i - p_\ell\|^2 / (2\sigma_\ell^2))},$$

where σ_j is the average distance to control j 's three nearest controls. The six nearest same-group controls form the shape-preserving ARAP graph. In plain terms, a player's splats follow that player's nearby controls, never a teammate's or the court's.

Shared Translation + Local Motion Trying to move every control independently made weakly observed controls lag or drift. By frame 10 that version reached 0.60 mask overlap (IoU), compared with 0.71 when each player first gets one shared translation and then small local offsets. The shared translation handles running across the court; the local offsets handle arms, legs, and other articulation. A rigid local MLS transform and the ARAP shape loss keep neighbouring splats from stretching apart.

At frame f , controls begin from the previous frame's velocity. We optimize one shared translation T_k per moving group and small local offsets δ_j for 50 Adam iterations, using four cameras per iteration. Their learning rates are 0.020 m and 0.003 m. Let $t_j = T_{g(j)} + \delta_j$ and $p'_j = p_j + t_j$. For Gaussian i , the fixed weights above blend nearby original controls p_j and moved controls p'_j to form a best local rotation R_i :

$$\bar{p}_i = \sum_j w_{ij} p_j, \quad \bar{q}_i = \sum_j w_{ij} p'_j, \quad M_i = \sum_j w_{ij} (p_j - \bar{p}_i)(p'_j - \bar{q}_i)^T.$$

With $M_i = USV^T$, set $R_i = VU^T$, flipping the last column of V if $\det(VU^T) < 0$. The final Gaussian position is $x'_i = R_i(x_i - \bar{p}_i) + \bar{q}_i$; its orientation and simple colour direction are rotated too. A fused CUDA operation performs this transform and its backward pass. The reported run does not separately optimize control rotations.

Every fifth frame, we render where each player is expected, prompt SAM2 with that box and centre, and assign overlapping pixels to the strongest rendered player. We triangulate the 2D mask centres from several cameras into a rough 3D player centre. This is a *motion anchor*: it gives the shared translation a reliable hint about where a player moved, especially after fast motion or occlusion. If enough cameras see the player, a robust Huber anchor of scale 0.3 m and weight 0.1 is used. Between those frames, no made-up mask or anchor is used. The deformation objective is

$$\mathcal{L} = \mathcal{L}_{\text{image}} + 0.3\mathcal{L}_{\text{BCE}}^{\text{instance}} + \mathcal{L}_{\text{ARAP}} + 0.01\mathcal{L}_{\text{accel}} + \mathcal{L}_{\text{anchor}},$$

Here the image term matches appearance, BCE matches each player's rendered mask, ARAP preserves local shape, acceleration discourages sudden jitter, and the anchor discourages global drift. Background controls stay frozen while the moving splats are rendered over a cached background.




Figure 1: CourtFlow-GS pipeline. The editable diagram source is distributed with the report as courtflow_gs_pipeline.drawio.

3.3 Focused refinement

MLS provides stable, smooth motion, but deliberately does not change a Gaussian’s colour, size, or amount of detail. Updating the full image after every video frame would be wasteful: players cover less than 1% of the court image, so most compute would repeatedly optimize the static floor. We therefore refine only the moving-player crops most frames, and use occasional full-image keyframes to correct drift and add a limited amount of new detail.

After every deformation, we save (“bake”) the moved Gaussian and control state for that frame. On a normal frame, all moving splats are projected into each training camera and surrounded by one box with 48 pixels of padding. The box is aligned to the rasterizer’s 16-pixel tiles, and the camera principal point is shifted so rendering the crop is exactly equivalent to rendering that part of the full image. We run 100 refinement iterations on two crops per step. No splats are added or removed here, so their control bindings stay stable.

Every tenth frame is a keyframe. It instead runs 500 full-image iterations, can add splats from iteration 50 every 50 iterations until 60% of the schedule, and limits growth to 5% per keyframe. It then reconnects each Gaussian to its eight nearest same-group controls and refreshes the cached background. Held-out cameras are never included in any training loss.

3.4 Pseudocode

Algorithm 1 summarizes the whole process. The helper names are intentionally literal: first build the frame-0 scene, then for each video frame estimate motion, bake that motion, and improve either player crops or a periodic full keyframe.

Algorithm 1 CourtFlow-GS online reconstruction

Require: Frame-0 images, synchronized training videos, fixed calibration

```
1:  $(state, controls) \leftarrow \text{BUILDCANONICALSCENE}(frame_0, calibration)$ 
2: for each  $f$  in the video sequence do
3:    $images \leftarrow \text{DECODEUNDISTORTNORMALIZE}(trainingVideos, f)$ 
4:   if  $f \bmod 5 = 0$  then
5:      $masks \leftarrow \text{REPAIRMULTIVIEWSAM2IDENTITIES}(images, state.dynamic)$ 
6:      $anchors \leftarrow \text{TRIANGULATEMASKCENTROIDS}(masks)$ 
7:   else
8:      $masks, anchors \leftarrow \emptyset, \emptyset$ 
9:   end if
10:   $offsets \leftarrow \text{OPTIMIZECONTROLS}(L_1 + \text{D-SSIM}, masks, anchors, \text{ARAP})$ 
11:   $state \leftarrow \text{BAKEFUSEDRIGIDMLS}(state, controls, offsets)$ 
12:  if  $f \bmod 10 \neq 0$  then
13:     $state \leftarrow \text{REFINEDYNAMICCROPS}(state, images, 100)$ 
14:  else
15:     $state \leftarrow \text{REFINEFULLIMAGES}(state, images, 500)$ 
16:     $controls \leftarrow \text{REBIND}(state, controls); \text{CACHEBACKGROUND}(state)$ 
17:  end if
18:   $\text{SAVEDIAGNOSTICS}(state)$ 
19: end for
```

4 Execution Walkthrough

The previous section explains *why* each part of the method exists. This section follows one complete run from the raw capture to the final videos: which program runs, what it reads, what it writes, and where time is spent. Everything is started by a single command,

```
python -m ring_init.e2e --scene basketball --dataset-dir <capture> \
    --config config.local.json --out-dir <out> --frames 0:700
```

where the capture directory holds the 36 videos (cameras/view_XXX.mp4) and the COLMAP calibration (calibration/). The local configuration only names the SAM2 checkpoint; all other settings are the defaults, and the exact resolved configuration is written to `e2e_config.json` in the output directory before any work starts. Every step writes a completion marker (`done.json`) next to its output, so an interrupted run resumes where it stopped and `--force` step recomputes a single step. The times below were measured on the RTX PRO 6000.

4.1 Preparing the input (17 s)

The capture adapter reads the calibration for all 36 cameras, decodes frame 0 of every video with `ffmpeg` (Lanczos resampling from 3840×2160 to the 1920×1080 working resolution), and removes lens distortion with `OpenCV` so every camera becomes an ideal pinhole camera. Every third view is copied as the 12 training images `cam_00...cam_11`; all 36 undistorted frame-0 images are kept for held-out evaluation. The adapter also writes the 12-camera and 36-camera calibration files used by the rest of the pipeline.

4.2 Stage A: building the frame-0 scene (≈ 9 min)

Stage A runs the steps of Table 3 in order. Each step caches its result on disk, and later steps only read those files, so the steps are independent programs joined by files.

Table 3: Stage A steps as executed, with RTX PRO 6000 wall-clock times.

Step	What runs	Writes	Time
calib	Loads the 12 training cameras and images; draws the camera layout.	calib/	2.1 s
match_full	MASt3R matches neighbouring ring cameras on whole frames (512 px).	3D points	33.4 s
floor	RANSAC fits the court plane and its extent; fixes metric scale.	floor plane	0.3 s
persons	Faster R-CNN person boxes, then SAM2 masks, in every training camera.	per-camera masks	11.7 s
instances	Carves the court volume with all masks; 3D blobs become player IDs; re-prompts SAM2 for missing views.	label images L_c	16.7 s
hulls	Occlusion-aware visual hull for every player.	voxel hulls	4.3 s
match_crops	MASt3R on masked player crops for nearby camera pairs.	player points	46.9 s
fuse	Merges crop matches and hull fill into coloured seed clouds.	seed .ply	5.6 s
train	Optimizes each player’s Gaussians, then the background with players composited in front.	Gaussian models	406.9 s
export	Copies the trained models into a canonical package with a manifest.	canonical/	<1 s
eval	Renders frame 0 for all 24 held-out cameras and scores it.	eval/	–

Training is the only substantial cost. The detector, SAM2, and MASt3R networks are freed from GPU memory once their steps finish, so none of them competes with Gaussian training for memory.

4.3 Stage B, part 1: decoding the sequence

Before tracking starts, Stage B prepares all per-frame data. Six worker processes each stream one camera’s video through ffmpeg and apply exactly the frame-0 preprocessing: the same resampling, the same precomputed undistortion map, and the same per-camera colour correction. The 12 training cameras are decoded at every frame (8400 images for 700 frames). The 24 held-out cameras are decoded only at every tenth frame, because they are used only for scoring. Frames are stored as lightly compressed PNG files, and already decoded frames are skipped on a rerun.

SAM2 then runs once per training camera in video mode. It is prompted with that camera’s frame-0 player labels from Stage A and follows every player through the whole sequence. These raw masks are not trusted as they are: SAM2 treats each camera separately and can swap two players in one camera but not another. Identities are therefore repaired in 3D during tracking (below). Finally, the ball is detected and triangulated at frame 0, tracked causally through the sequence with a gravity and bounce model, and added to the canonical scene as its own small group of Gaussians.

4.4 Stage B, part 2: the per-frame loop

The canonical Gaussians are loaded onto the GPU once and stay there for the whole sequence. Control points are placed on every group (1024 per player, 16 on the ball, 1000 on the background), and each Gaussian is bound to its eight nearest same-group controls. Then, for every frame f :

1. **Load.** A background thread has already read frame f ’s 12 training images from disk while the GPU was busy with frame $f - 1$. The images are downsampled to the tracking resolution (960×540) and uploaded; the full-resolution copies are kept for refinement.
2. **Masks and anchors (every fifth frame).** The SAM2 masks for frame f are read, their identities are rebuilt in 3D exactly as in Stage A’s `instances` step, and the repaired masks are matched to the tracked

players by position and jersey colour. Triangulating each player’s mask centre gives a motion anchor. On other frames, no masks or anchors are used.

3. **Deform.** Fifty Adam iterations update each player’s shared translation and local control offsets; the fused CUDA MLS operation moves the Gaussians. Background controls stay frozen, and the moving players are rendered over a cached background image.
4. **Improve detail.** On nine frames out of ten, 100 iterations refine only the crops around moving players. Every tenth frame is a keyframe. First, a health check compares each player’s rendered footprint with the repaired masks; a player whose overlap has fallen below 0.10 IoU is re-initialized from its canonical shape at its anchor. Then 500 full-image iterations run with limited densification, the Gaussians are re-bound to their controls, and the background cache is refreshed.
5. **Record.** The current state is rendered for held-out views 13 and 26 and for the orbit camera and streamed to ffmpeg, so each video frame shows the scene at that instant. Every tenth frame, and at the last frame, all 24 held-out cameras are rendered and scored (PSNR, SSIM, LPIPS, and player-region PSNR); the training views also report mask IoU. Every 25 frames, the full Gaussian state is saved as a PLY checkpoint and the control trajectories and metrics are flushed to disk.

After the last frame, the final state is saved and per-stage timings and quality statistics are summarized in `summary.json`. The held-out view-13 video and the orbit video are then published to `final_renders/` with an `animated_` prefix, marking them as renders of the moving scene rather than of one baked checkpoint.

4.5 What the run leaves on disk

The output directory keeps every intermediate product, so any step can be inspected or rerun:

```
<out>/prepared_input/basketball/  undistorted frame-0 images and calibration
<out>/basketball/
  e2e_config.json, timings.json  resolved settings and step times
  persons/ instances/ hulls/ ... Stage A intermediate results
  canonical/                    editable frame-0 Gaussian scene + manifest
  eval/                         frame-0 held-out scores and images
  frames/ eval_frames/          decoded training and held-out frames
  video_masks/ ball/           SAM2 masks and the ball trajectory
  stage_b_v2/
    metrics.json summary.json  per-frame metrics and the run summary
    control_positions.npz      control-point trajectories for all frames
    ply/frame_NNNNNN/         Gaussian checkpoints every 25 frames
    videos/ final_renders/     held-out and orbit videos
```

5 Results

The final diagnostic reconstructs 300 frames. PSNR measures image similarity in decibels (higher is better); person PSNR measures it only on player regions; IoU measures how well predicted and true player masks overlap (1 is perfect). The diagnostic’s held-out result is the important generalization measurement because those cameras were never trained on.

Table 4: Quality summary. Temporal held-out samples coincide with keyframes.

Method	Full PSNR	Person PSNR	Person IoU	Measurement
Global 3DGS baseline	17.50	—	—	frame-0 held-out
Stage A canonical	19.32	18.34	0.788	frame-0 held-out
Stage B v1	18.29	15.00	0.70→0.30	keyframe held-out mean
Stage B v2 (old)	18.87	19.62	0.71→0.39	keyframe held-out mean
Focused refinement, frames 1–20	22.42	26.76	0.746	training-view mean
Final 300-frame diagnostic	20.07	22.44	—	held-out mean



Ground truth, view 13, frame 200



Earlier Stage B v2



CourtFlow-GS diagnostic

Figure 2: Same-camera comparison at held-out view 13, frame 200. Replace or add comparison rows as desired.

5.1 Why these design choices were selected

The method was not chosen by adding complexity for its own sake. Table 5 records the practical failures we observed and the smallest change that addressed each one. These are useful trade-offs: some choices improve quality and speed at the same time by avoiding work on irrelevant pixels, while others deliberately spend more computation only where player detail needs it.

Table 5: Design choices and the observation behind each one.

Choice	Why it was selected
3D-first identities	Matching 2D centres and colours independently fragmented 68 masks into 60 apparent people. Building rough 3D components first gives the same player one ID across cameras.
Occlusion-aware hull	A strict silhouette intersection removed hidden limbs. Allowing arbitrary misses inflated a hull by about $2.5\times$. Measured nearer depth supplies the middle ground.
Masked MAST3R crops	Full-frame descriptors reduce a small player to about 6×25 pixels and produced star-shaped rays. Crops spend the fixed descriptor resolution on the player instead of the court.
Same-group MLS controls	A global control graph lets nearby teammates pull each other. Fixed same-player control neighbours keep motions separate and also make deformation cheap.
Shared translation plus local offsets	Moving every control alone reached 0.60 IoU by frame 10; shared translation plus local articulation reached 0.71 IoU.
Crop refinement plus keyframes	Full-image updates mostly optimize static court pixels. Player crops improve moving detail cheaply; full keyframes occasionally correct accumulated error and add splats.

5.2 Evaluation protocol and deliverables

The 12 training cameras are never used as novel-view evidence. At every tenth frame and at the final frame, the tracker renders all 24 held-out cameras and compares them to their synchronized real images. The current report gives full-frame PSNR and player-region PSNR; player-region IoU checks whether the rendered player footprint agrees with the independent mask. The final RTX PRO 6000 run will report the mean and standard deviation over the held-out camera/frame pairs, together with SSIM and LPIPS, and will include a grid of easy, crowded, and occluded held-out views. This makes it clear that quality is measured from cameras the model did not optimize against.

The release command produces two animated deliverables during tracking: a held-out camera video and an orbit video. These renders use the state from each corresponding video frame. This matters because one baked PLY checkpoint is a snapshot at one instant; repeating the final snapshot would make a static movie, not an animation. The editable canonical model, controls, and calibration are kept for re-training or editing. The final source scene is about 150 MB on the RTX 3080; a naive per-frame PLY bake would be about 42 GB for 300 frames, so a deployable animated package should use packed, quantized frame chunks rather than text PLY files.

6 Runtime

Table 6: Runtime on both GPUs (preliminary for the PRO 6000). Working resolution: RTX 3080 1874×1054 with tracking renders at 937×527 ; RTX PRO 6000 1920×1080 with tracking renders at 960×540 (5% more pixels, so the PRO 6000 speed-ups are slightly understated). Stage B rows are steady-state means over the same frames 2–29 (28 frames, two keyframes) with the identical production configuration, including 3D identity repair, on both GPUs. Frame 1 is excluded because the 3080 compiled its CUDA kernels there (159 s); the PRO 6000 compiles them beforehand. The raster row renders a baked checkpoint at the stated resolution.

Component	RTX 3080 12 GB	RTX PRO 6000	Speed-up
Stage A frame-0 canonical build	1896 s	528 s	$3.6\times$
MLS deformation/control optimization	2.75 s/frame	1.28 s/frame	$2.1\times$
of which fused MLS forward	0.072 s/frame	0.029 s/frame	$2.5\times$
Focused crop refinement (100 iterations)	3.04 s/non-keyframe	0.92 s/non-keyframe	$3.3\times$
500-step full keyframe	23.02 s/keyframe	5.75 s/keyframe	$4.0\times$
Mask refresh with identity repair, amortized	1.83 s/frame	1.02 s/frame	$1.8\times$
Stage B total	9.05 s/frame (0.110 FPS)	3.57 s/frame (0.280 FPS)	$2.5\times$
GPU SH colour + baked raster, 937×527	4.36 ms (229 FPS)	[measure on RTX PRO 6000]	

The fused MLS/SVD calculation is only a small part of tracking time. Rendering and differentiating through the rasterizer dominate optimization. Once motion is baked, GPU rendering is fast; the current video exporter is instead limited by copying frames from Python to ffmpeg.

The PRO 6000 gains most where the work is large, dense Gaussian optimization: keyframes are $4.0\times$ faster and crop refinement $3.3\times$. Control-point optimization gains less ($2.1\times$); each of its 50 iterations renders only four cropped, half-resolution cameras, a workload too small to occupy the larger GPU fully. Mask refresh gains least ($1.8\times$); part of the 3D identity repair (mask decoding, assignment, and relabelling) runs on the CPU, which the new GPU does not accelerate.

These Stage B values describe the early part of the sequence. Later, crowded frames cost more: over all 300 frames of the same configuration the RTX 3080 averaged 10.54 s/frame, against 9.05 s/frame over frames 2–29. The complete 700-frame PRO 6000 run will replace these preliminary values with full-sequence averages, total wall-clock time, peak VRAM, and baked rendering time at 1080p and 4K.

7 Limitations and next steps

The main remaining weakness is identity tracking in crowded late-game plays. Independent SAM2 masks can swap people or disappear. In the previous diagnostic, training IoU falls from 0.789 near the start to 0.001 at frame 299; held-out person PSNR falls from 23.86 to 20.59 dB. The production configuration now repairs SAM2 identities in 3D at mask frames and performs conservative keyframe recovery when a repaired mask shows that a player was lost. The 300-frame repair run tests whether that change removes the observed failure rather than merely moving it later in the sequence. Other useful work is a smaller or level-of-detail version of the 2.01M-Gaussian final state, unbiased inter-keyframe held-out evaluation, and direct GPU-to-NVENC video output.

References

- [1] J. Luiten, G. Kopanas, B. Leibe, D. Ramanan. Dynamic 3D Gaussians: Tracking by Persistent Dynamic View Synthesis. 3DV 2024. <https://arxiv.org/abs/2308.09713>
- [2] J. Sun et al. 3DGStream: On-the-Fly Training of 3D Gaussians for Efficient Streaming of Photo-Realistic Free-Viewpoint Videos. CVPR 2024. <https://arxiv.org/abs/2403.01444>
- [3] S. Girish et al. QUEEN: QUantized Efficient ENcoding of Dynamic Gaussians for Streaming Free-viewpoint Videos. NeurIPS 2024. <https://arxiv.org/abs/2412.04469>
- [4] J. Liang et al. ClipGStream: Clip-Stream Gaussian Splatting for Any Length and Any Motion Multi-View Dynamic Scene Reconstruction. CVPR 2026. <https://arxiv.org/abs/2604.13746>
- [5] Y.-H. Huang et al. SC-GS: Sparse-Controlled Gaussian Splatting for Editable Dynamic Scenes. CVPR 2024. <https://arxiv.org/abs/2312.14937>
- [6] H. Park, G. Ryu, W. Kim. DropGaussian: Structural Regularization for Sparse-view Gaussian Splatting. CVPR 2025. <https://arxiv.org/abs/2504.00773>